

Sharpening the Brain

Training adjusts the network through repeated feedback



AI

NAMASTE AI - SEASON 1

Episode 07

Sharpening the Brain

This episode moves from an already-trained transformer to the process that makes it useful, explaining parameters, forward prediction, loss, backpropagation, gradients, optimizers, self-supervised targets, generalization, overfitting, and the repeated training loop.

02:01:38

WebVTT transcript only

In this episode

01 What learning means for a neural network**02** Parameters, training data, and prediction targets**03** Forward pass and loss**04** Backpropagation, gradients, and optimizers**05** Gradient descent and learning rate**06** The complete self-supervised training loop**07** Training, inference, generalization, and overfitting**08** Distributed training, learned embeddings, and understanding

From a trained model to an untrained one

00:00 - 06:43

The previous episode followed data through a transformer whose parameters were already useful. Given `The pizza is ready`, it could assign a high score to a continuation such as `to`. This episode removes that assumption.

Start instead with an untrained or poorly trained neural network:

Trained

The sky is ... may continue with blue .

Untrained

The sky is ... may continue with potato , banana , cool , magic , or another random-looking token.

The instructor compares the untrained model with a newborn child for whom the world initially appears like gibberish. The comparison is intentionally broad; the episode's real question is numerical: how does a network move from random output to a better prediction?

Q: What does “learning” mean inside this model?

A: Learning means repeatedly adjusting the model's parameters so its future predictions become better for the chosen objective. Here, the objective is next-token prediction.

Parameters: the adjustable numbers inside the network

06:43 – 19:34

A neural network contains an enormous collection of adjustable numbers called **parameters**. Those numbers participate throughout the transformer's mathematics: embedding lookup, attention, feed-forward processing, matrix multiplication, dot products, addition, division, and output scoring.

The compact definition is:

learning = adjusting parameters so future predictions become better

Three analogies make small numerical updates easier to picture.

A DJ controller full of knobs

Imagine a large controller with many knobs. Every knob is one adjustable value. Turning any one of them slightly can make the output better or worse. Training gradually moves the full collection toward useful positions.

An old radio

An old radio must be tuned to a precise frequency. A small turn can reduce noise and make the channel clearer. Now imagine thousands, millions, or billions of knobs being tuned automatically.

A guitar

Later, the lecture compares training with tuning a guitar and inference with playing it. Training changes the instrument; inference uses its tuned state.

An illustrative parameter might move from 2.5 to 2.4 , then to 2.2 , with each small change producing a slightly better result. These are not values from an actual optimization trace.

The scale discussion names GPT-3 as having **175 billion parameters**, also expressed as **17,500 crore**. The instructor then says later GPT-4/GPT-5 counts were not publicly revealed, but a later drawing uses **1.8 trillion parameters** without identifying its source.

Treat **175 billion** as the instructor's stated GPT-3 figure. The **1.8 trillion** number is an ambiguous scale example; the transcript does not establish which model it refers to.

The purpose of these large numbers is not memorization. It is to show that training makes tiny coordinated adjustments across a huge numerical system.

Which values are parameters?

19:34 – 23:40

The instructor revisits the previous transformer architecture and identifies several learned value families:

- token-embedding table values;
- Q, K, and V attention weights;
- output or projection weights;
- bias values;
- layer-normalization gamma and beta;
- MLP or feed-forward projection weights and biases.

In an untrained network, these values begin as random numbers. Training modifies them.

At **20:33**, the transcript says “Token ID is embeddings.” The surrounding course consistently distinguishes token IDs from embeddings, and the next line identifies embeddings as parameters. These notes retain that established distinction.

Do parameters store knowledge?

23:40 – 27:46

This is presented as both an interview question and a philosophical distinction. The instructor gives a deliberately nuanced answer:

- parameters do not contain a database-like paragraph or literal memory record;
- they are numerical values rather than a lookup table of facts;
- learned patterns are distributed across those values;
- when input passes through those patterns, the network can produce a response that appears knowledgeable;
- the instructor calls parameters “**knowledge enablers.**”

Earlier wording says that “patterns and knowledge” are encoded across parameters, while this passage says they do not “store knowledge or memory.” The intended distinction is between distributed learned patterns and database-like storage, but the lecture does not reduce that distinction to one perfectly consistent sentence.

INTERVIEW NOTE

In the lecture’s framing, parameters are not a database of memorized paragraphs. They are learned numbers whose distributed patterns shape the model’s output.

Training data is not the same as parameters

23:40 – 27:46

Training data is the external information supplied to drive learning. The instructor imagines large collections of public internet content, including articles, blogs, Wikipedia, social-media material, images, photos, documents, code, history, geography, literature, poetry, songs, and language examples.

The data is tokenized and represented numerically before passing through the network. Parameters live inside the network and are updated; training data supplies the examples that cause those updates.

The terms terabytes and petabytes communicate scale. The instructor later adds an important qualification: web data should be crawled, cleaned, selected, and filtered rather than fed to the model indiscriminately.

Phrases such as “all possible public information on the internet” are simplifying lecture language, not a documented inventory for a particular model.

[Transcript unclear] Captions around 26:37–27:09 are badly garbled. The data-versus-parameters distinction resumes clearly after that interval.

The forward pass and a known target

27:46 – 33:32

A **forward pass** is the prediction-producing trip through the network:

- 1 Take one input sample.
- 2 Convert it into the model’s numerical representation.
- 3 Process it through the transformer layers.
- 4 Produce scores or probabilities for the next token.

The source text itself supplies a correct next token. From:

The sky is blue

training can construct:

Role	Text/token
input sample	The sky is
known target	blue
untrained prediction	perhaps banana

One mock-up includes values such as `beautiful = 0.55` , `blue = 20` , `banana = 85` , `green = 61` , and `best = 0.15` . A separate example uses `blue = 2%` and `banana = 80%` . These scales are inconsistent and should not be merged into one normalized table. Their shared message is simply that an untrained model may score the wrong token above the known target.

Training should move the system toward a larger score for `blue` and a smaller score for `banana` .

[Transcript unclear] Around `30:32-31:08` , the lecture briefly starts `born and raised in New ...` with target `zealand` , then switches back to `The sky is ... blue` while still describing the same prediction process.

At `32:51` , a brief sentence says `banana` should receive a “bigger” number; the surrounding explanation consistently says banana should decrease while blue increases.

Loss: measuring how bad the prediction is

33:32 – 40:01

The lecture defines **loss** as a numerical measure of how bad a prediction is. A **loss function** turns prediction quality into a number.

The intuitive cases are:

Small loss

The known target `blue` receives a high probability.

High loss

`blue` receives `2%` while the wrong token `banana` receives `80%` .

No exact loss equation is spoken. **[Exact loss function not available]** Captions around `34:01-34:27` , `35:51-35:57` , and `36:48-37:18` interrupt the explanation, but the definition and loop are repeated clearly elsewhere.

The first training cycle now looks like this:

`predict → compare with target → calculate loss → use the loss to improve parameters`

The simplified examples sometimes treat a perfect prediction as loss 0 with no further update. Later, the instructor says a real model's loss can never be exactly zero, then immediately says he should not declare that impossible. The stable learner-facing idea is that training tries to reduce loss.

Backpropagation: trace responsibility backward

40:01 – 47:50

With billions of parameters, a single loss value is not enough. Which values contributed to the bad result, and how should each one change?

Backpropagation works backward from the output error through the transformer layers. If the forward path is:

Layer A → Layer B → Layer C → prediction → loss

then the backward investigation follows:

loss → Layer C → Layer B → Layer A

The lecture's detective analogy asks which layer and which parameter was the "culprit." The more precise explanation is that backpropagation calculates how changing each parameter would affect the loss. The investigation reaches embedding values, attention weights, feed-forward weights, and biases.

Gradients are sensitivities

The calculated sensitivity of loss to a parameter is called a **gradient**. In the instructor's language, a gradient provides:

- a direction in which a parameter should move;
- an indication of how sensitive the loss is to that parameter.

For example, begin with a parameter at 0.40. If increasing it would make the loss larger, move downward-perhaps to 0.39 or 0.38 rather than taking a large jump.

The instructor uses *weight* and *parameter* interchangeably in this passage, although the earlier parameter list includes embeddings and biases as well as weights. The safe distinction is that weights are among the parameters; the lecture sometimes uses the shorter word loosely.

The lecture also loosely associates larger loss with larger gradient. Keep that only as its intuition, not as an equation.

Backpropagation versus the optimizer

47:15 – 50:05

The instructor draws a clear division of work:

Backpropagation

Computes gradients: direction and sensitivity for the parameters.

Optimizer

Uses those gradients to update the parameter values.

In this explanation, backpropagation diagnoses; the optimizer changes the numbers. Later, around **01:50:45–01:51:05**, the lecture says that “during backpropagation” embedding values are adjusted. These notes retain the earlier, explicitly stated separation and treat the later phrase as loose shorthand for the training process.

Gradient descent: repeated movement toward lower loss

47:50 – 56:49

The lecture defines **gradient descent** as an iterative optimization process that tries to minimize a function such as loss by adjusting parameters in the opposite direction of the gradient.

It needs three ingredients:

- the loss function;
- gradients;
- a learning rate.

Why the opposite direction? If increasing a parameter increases the loss, the optimizer should move that parameter downward.

The **learning rate** controls the step size. A value at **0.4** might move to **0.39** rather than leaping straight to **0.2**.

The foggy-mountain analogy

Imagine standing on a mountain in thick fog. The goal is to reach the lowest point, but the full route is invisible.

- 1 Inspect the local slope.
- 2 Take a small step downhill.
- 3 Inspect the slope again from the new position.
- 4 Repeat until reaching a low point.

Mountain analogy	Training concept
height	loss
downhill direction or slope	gradient direction
size of each step	learning rate
repeated small steps	iterative parameter updates
bottom	minimized loss

Gradient descent

The lecture uses a foggy-mountain analogy: follow local slope information toward lower loss



In the foggy-mountain analogy, the local slope supplies a direction and the learning rate controls each step toward lower loss.

Across repeated examples, `The sky is ...` might first produce banana, then potato, then beautiful, and eventually blue. Every poor result supplies information for another small update.

Batch, stochastic, and mini-batch gradient descent

56:49 – 59:22

The variants differ in how much data contributes to one update:

Approach	Data used per update	Lecture trade-off
Batch gradient descent	Entire dataset	Stable, but slow on a huge dataset.
Stochastic gradient descent (SGD)	One data point	Fast per step, but noisy and repeated often.
Mini-batch gradient descent	A small subset or batch	A balance between speed and stability-the lecture calls it the “best of both worlds.”

The instructor says he is unsure how to pronounce *stochastic*, but the term and comparison are clear.

Optimization can still get stuck

59:22 – 01:01:33

Three challenges extend the mountain analogy:

- **learning rate too small:** progress can take extremely long;
- **learning rate too large:** the optimizer can overshoot the minimum and fail to converge;
- **local minimum, saddle point, or flat region:** the local slope may provide little guidance toward a better path.

Picture a plateau hidden by fog. The surface looks flat, so local information does not clearly reveal where a deeper descent begins.

The complete training loop


01:01:33 – 01:08:54

All of the pieces now form one repeated cycle:

- 1 **Sample:** take a training example.
- 2 **Forward pass:** process it through the transformer.
- 3 **Prediction:** produce next-token scores.
- 4 **Target comparison:** compare the prediction with the actual next token in the source text.
- 5 **Loss:** measure the error numerically.
- 6 **Backpropagation:** compute gradients backward through the network.
- 7 **Optimizer:** update parameters using the gradients and learning rate.
- 8 **Repeat:** continue with more samples or batches using a slightly improved model.

The complete training loop

A prediction is evaluated, responsibility flows backward, and parameters are updated



Training repeatedly evaluates a prediction, traces responsibility backward, updates the parameters, and begins again with another sample or batch.

The instructor calls a forward pass one attempt to predict a token. In a later sentence he also uses “one pass” for the broader prediction/loss/backprop/update sequence. To keep the concepts distinct, these notes use **forward pass** for the prediction-producing computation and **training step** for the complete update cycle.

A human-learning analogy for repeated patterns

01:08:54 – 01:20:00

Familiar phrases become easy for people to complete:

- “Honesty is the best ...” → **policy** ;
- “The sun rises in the ...” → **east** ;
- “Roses are ...” → **red** ;
- after repeatedly hearing “Namaste AI is beautiful,” the phrase “Namaste AI is ...” invites **beautiful** .

The instructor uses these completions to picture a progression in model learning:

1. random or gibberish output;
2. familiar word-sequence patterns;
3. grammar and singular/plural patterns;
4. contextual relationships through attention;
5. longer-range relationships;
6. increasingly useful language, code, and domain output.

Repeated code syntax provides a software-engineering analogy. After seeing `if` statements many times, a learner begins to expect their brackets or curly braces; the model similarly learns recurring token patterns.

The instructor repeatedly qualifies this comparison. The human brain is far more sophisticated, and the link between neural-network parameter changes and biological-neuron changes is only a high-level philosophical analogy.

[Transcript unclear] Several grammar examples involving `I`, `we`, `what`, `you`, and `a/an` are fragmented around `01:16:45-01:17:44`. Their shared point-learning grammatical co-occurrence patterns-is recoverable, but the individual examples are not.

Compute, GPUs, and training scale

01:20:00 – 01:23:49

The full loop-forward prediction, loss, backpropagation, gradients, and optimization-must be repeated over enormous datasets. The instructor connects that repetition with substantial compute, power, GPUs, storage and data handling, optimized software, and distributed infrastructure.

He contrasts humans, who tire and read relatively few books, with machines that can keep processing while compute and power remain available. Comments about companies competing for GPUs, Nvidia's contemporary value, and named industry figures are time-sensitive instructor observations rather than verified facts for these notes.

The lecture's broad wording that more data produces a better model is later balanced by its own requirement to clean and filter the data. The notes preserve both: scale matters in the explanation, and indiscriminate data is not presented as the target.

Why next-token learning is self-supervised

01:23:49 – 01:27:57

The instructor first says "supervised self-learning," then corrects himself to **self-supervised learning**.

The target can come directly from the text:

- 1 Start with `The sky is blue`.
- 2 Use `The sky is` as the input.
- 3 Hide/use `blue` as the target next token.
- 4 Predict, compare with the target, calculate loss, and update.

The compact loop is:

`predict → evaluate → fix`

No person has to label the correct next word for every sentence; the sequence supplies it. The instructor contrasts this with a cat/dog image example in which humans attach labels to images. That contrast is a broad teaching history, not a complete taxonomy of supervised and self-supervised learning.

Before any of this, web data must still be crawled, cleaned, selected, and filtered so the system is not trained indiscriminately on fake or unwanted content.

Training vocabulary: sample to epoch

01:27:57 – 01:33:58

Term	Lecture definition
sample / training example	One input used to train or test the network; it may be a word, sentence, paragraph, or larger piece.
dataset	The collection of all training examples.
batch	A group of examples processed together.
training step	One optimizer update.
epoch	One pass through the training dataset.
context window	The bounded input/context region the model processes.

Batches allow the system to process groups in parallel instead of handling every sentence strictly one at a time. The instructor says loss is averaged over the batch before the parameter update.

The context-window mechanics during training are described only briefly as the region filling or shifting while prediction proceeds.

Training versus inference

01:33:58 – 01:38:02

Both training and inference use the model’s forward architecture, but they continue differently:

Training	Inference
Runs a forward pass.	Runs a forward pass.
Compares the prediction with a target and calculates loss.	Uses the prediction as output.
Runs backpropagation and computes gradients.	Does not update parameters in this explanation.
Optimizer updates the parameters.	Uses already-tuned parameters.
Changes the model.	Uses the model.
More compute-intensive and slower.	Comparatively lighter and faster per generated token.

Training and inference

The same model participates in two different processes

Training

- input plus known target
- prediction and loss
- backpropagation
- parameters are updated

Inference

- new input or prompt
- forward computation
- output is generated
- parameters stay fixed

Training tunes the model by continuing from prediction to loss and parameter updates; inference uses the tuned parameters to produce output.

The guitar analogy returns: training tunes the strings; inference plays the tuned instrument.

The instructor first says inference does not need GPUs, then immediately says it still needs compute, servers, cooling, and power. The consistent comparison is relative: training consumes far more compute than inference.

An environmental aside says servers use water for cooling and “may” contribute to global warming, but the instructor qualifies it with “I guess.” It remains an attributed aside rather than a verified conclusion.

Generalization: use the pattern on something new

01:38:02 – 01:41:42

Generalization means performing well on examples the model did not directly memorize.

Training-like statements include:

- “The sky is blue.”
- “The ocean is blue.”
- “Grass is green.”
- “Leaves are green.”

Now give the unseen prompt:

The clear afternoon sky looked ...

If the model produces **blue** even though that exact sentence was not in its training examples, the instructor calls that generalization. “Good learning” finds patterns that transfer beyond exact repetitions.

Overfitting: good on training examples, weak on new ones

01:41:42 – 01:45:03

Overfitting is the opposite failure mode in this explanation. Suppose a dataset repeats **The sky is blue** 10,000 times. The model may become excellent at completing **The sky is ...** with **blue**, yet fail on:

On a clear summer afternoon, the sky appeared ...

Generalization	Overfitting
Works on new samples, not only memorized ones.	Performs very well on repeated/training samples but poorly on new ones.

The instructor highlights both terms as interview vocabulary worth being able to explain.

Training is also a distributed-systems problem

01:45:03 – 01:47:43

Large-scale training cannot consume internet-scale data strictly one item at a time. Work is divided and processed across hardware. The instructor therefore calls it both a machine-learning problem and a distributed-systems problem.

The factors he names are:

- more online data;
- transformer and other algorithms;
- improving GPUs and hardware;
- better GPU software;
- distributed-compute capability.

An **H100** GPU is mentioned as a powerful machine that can be rented online for private-model work. No provider, link, price, or specification is supplied.

The broad statement that “all training happens in parallel” is followed by a garbled caveat. **[Transcript unclear]** Do not strengthen it into a universal claim about every training architecture.

— Embeddings are parameters that learn too

01:47:43 – 01:53:42

The episode now closes the loop with Episode 05. The initial vectors for **king** and **queen** do not begin with a meaningful relationship; their embedding values are random parameters.

Across training:

1. predictions create loss;
2. backpropagation computes gradients for embedding values along with other parameters;
3. the optimizer updates those values;
4. repeated usage patterns make the representations more useful and related;
5. no human manually drags **king** and **queen** next to each other.

Embeddings are called **learned** because better representations help the network make better predictions.

The instructor later uses “backpropagation adjusts embeddings,” but this should be read alongside the earlier distinction: backpropagation computes gradients, while the optimizer applies parameter updates.

— Does prediction count as understanding?

01:53:42 – 01:59:43

The final technical transition becomes an explicitly personal philosophical discussion:

- A model learns patterns, but does it understand meaning?
- Is next-token prediction the same as human understanding?
- If the model responds appropriately to “I am sad,” does it understand sadness?
- Could a machine ever have feelings, awareness, or consciousness?
- What definition of *understanding* would settle the disagreement?

The instructor leans toward saying that mathematical token prediction is not human-like understanding. A word may even be split into several tokens, and the model operates over numerical representations. But he also presents the opposing position: if a system responds as though it understands, some people may count that behavior as understanding.

The episode does not resolve the debate. It says the definition is unclear and leaves the question open.

Examples of possible splits for `king`, `JavaScript`, and `impossible` are illustrative only. The lecture does not provide exact tokenizer output, so no token boundaries are invented here.

Closing thought

01:59:43 – 02:01:38

The episode closes the first training picture without claiming that every detail is finished. More topics remain in the course. The instructor asks learners to share the course and ends with “Namaste.”

Chapter summary

For this neural network, learning means changing parameters so future next-token predictions improve. Training data supplies samples and known targets; parameters are the internal learned numbers distributed through embeddings, attention, normalization, and feed-forward layers.

A forward pass produces a prediction. A loss function measures how far that prediction is from the target. Backpropagation works backward to compute gradients, and an optimizer such as the gradient-descent process described here uses those gradients and a learning rate to update the parameters. The full sample-to-update cycle repeats across batches and epochs.

Next-token pre-training is called self-supervised because the text itself supplies the target. Training changes the model; inference uses its tuned state. Useful learning should generalize to unseen examples rather than overfit exact repetitions. At large scale, the process also becomes a compute, GPU, software, and distributed-systems challenge. Even the meaningful embedding relationships from Episode 05 arise through this update loop.

Key takeaways

- Learning means adjusting parameters to reduce future prediction error.
- Parameters include embeddings, attention weights, biases, normalization values, and MLP projections.
- Training data supplies examples and targets; parameters are the values changed inside the network.
- A forward pass predicts, and loss measures how bad the prediction is.
- Backpropagation computes gradients; the optimizer applies updates.
- Gradient descent uses repeated steps, with the learning rate controlling step size.
- Self-supervised next-token learning gets its target from the sequence itself.
- Training changes parameters; inference uses trained parameters.
- Generalization transfers a learned pattern to unseen examples, while overfitting does not.
- Embedding relationships are learned because useful representations help reduce loss.

Revision questions

1. What is the difference between trained and untrained output for `The sky is ?`
2. State the episode's definition of learning.
3. How do the DJ-knob, radio, and guitar analogies each explain parameters?
4. Which parts of the Episode 06 transformer are identified as parameters?
5. How does the instructor distinguish distributed learned patterns from database-like memory?
6. What different roles do training data and parameters play?
7. What happens during a forward pass?
8. How does `blue` become a self-supervised target from `The sky is blue`?
9. What does loss measure?
10. Why is `blue = 2%` and `banana = 80%` a high-loss output?
11. What does backpropagation compute as it travels backward through the layers?
12. What is a gradient in the lecture's terminology?
13. What is the difference between backpropagation and the optimizer?
14. Define gradient descent in the wording of the episode.
15. What does the learning rate control?
16. Map height, slope, step size, and the mountain bottom to their training concepts.
17. Compare batch, stochastic, and mini-batch gradient descent.
18. What happens when the learning rate is too small or too large?
19. How can a flat or saddle region make optimization difficult?
20. Reconstruct the complete training loop in the correct order.
21. Define sample, dataset, batch, training step, epoch, and context window.
22. Why is next-token pre-training called self-supervised in this episode?
23. How do training and inference use the same forward architecture differently?

24. Explain generalization using the unseen afternoon-sky sentence.
25. Explain overfitting using the repeated `The sky is blue` example.
26. Why does the instructor call large-model training a distributed-systems problem?
27. How do initially random `king` and `queen` embeddings become related?
28. What two opposing views about machine understanding does the instructor present?